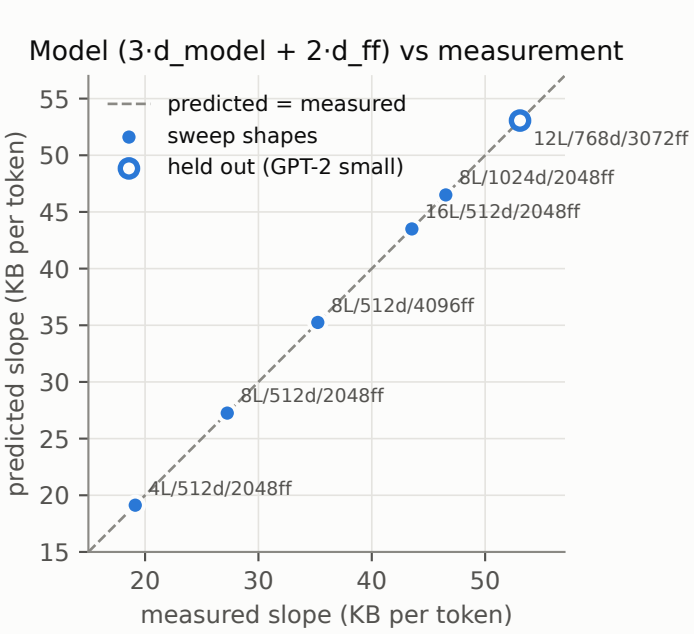
Model (3·d_model + 2·d_ff) vs measurement
predicted = measured
sweep shapes
held out (GPT-2 small)
12L/768d/3072ff
8L/1024d/2048ff
16L/512d/2048ff
8L/512d/4096ff
8L/512d/2048ff
4L/512d/2048ff
predicted slope (KB per token)
measured slope (KB per token)
15
20
25
30
35
40
45
50
55
20
30
40
50

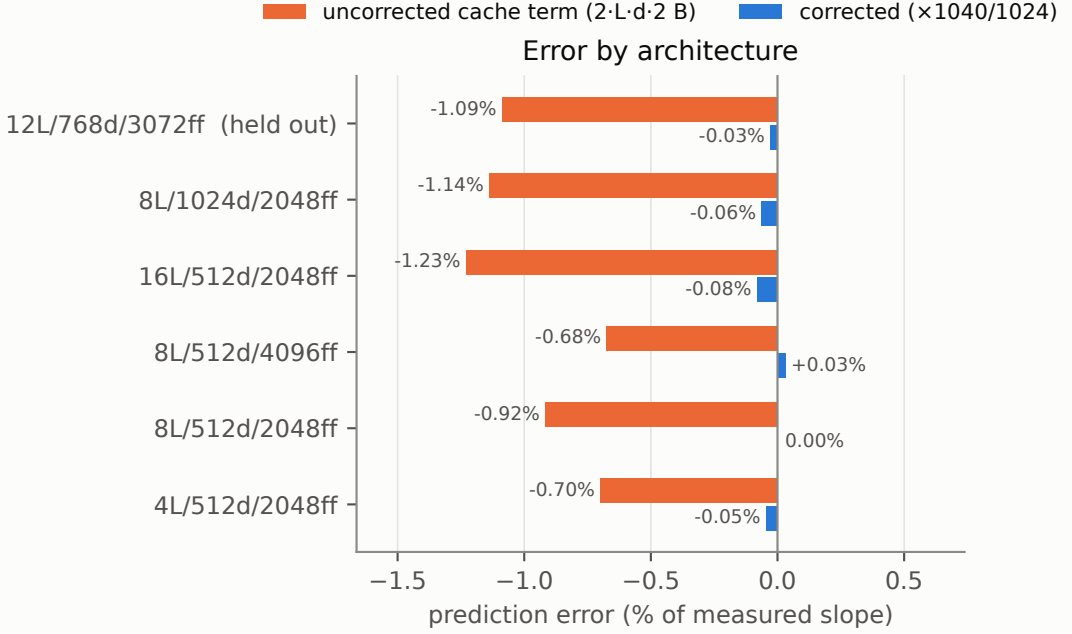
uncorrected cache term (2·L·d·2 B)
corrected (×1040/1024)
Error by architecture
12L/768d/3072ff (held out)
8L/1024d/2048ff
16L/512d/2048ff
8L/512d/4096ff
8L/512d/2048ff
4L/512d/2048ff
-1.09%
-0.03%
-1.14%
-0.06%
-1.23%
-0.08%
-0.68%
+0.03%
-0.92%
0.00%
-0.70%
-0.05%
−1.5
−1.0
−0.5
0.0
0.5
prediction error (% of measured slope)