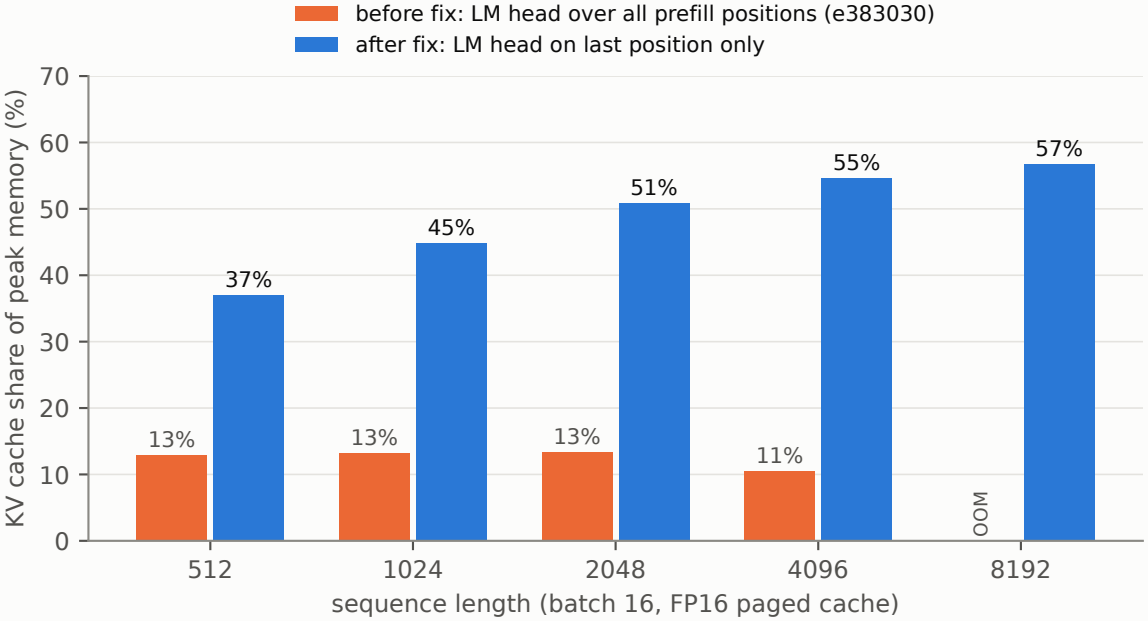
before fix: LM head over all prefill positions (e383030)
after fix: LM head on last position only
KV cache share of peak memory (%)
70
60
50
40
30
20
10
0
13%
37%
13%
45%
13%
51%
11%
55%
OOM
57%
512
1024
2048
4096
8192
sequence length (batch 16, FP16 paged cache)